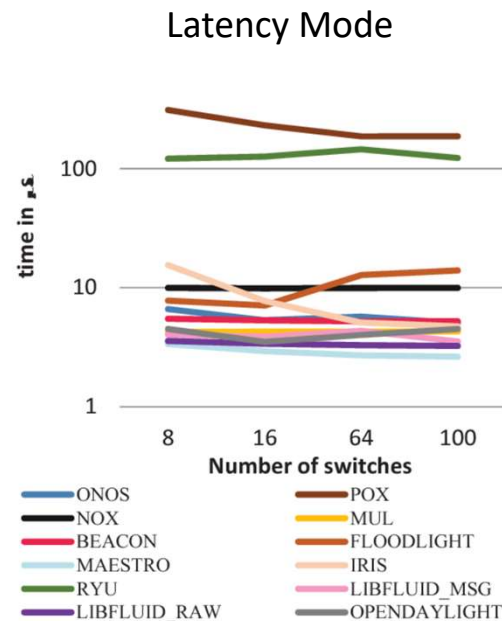
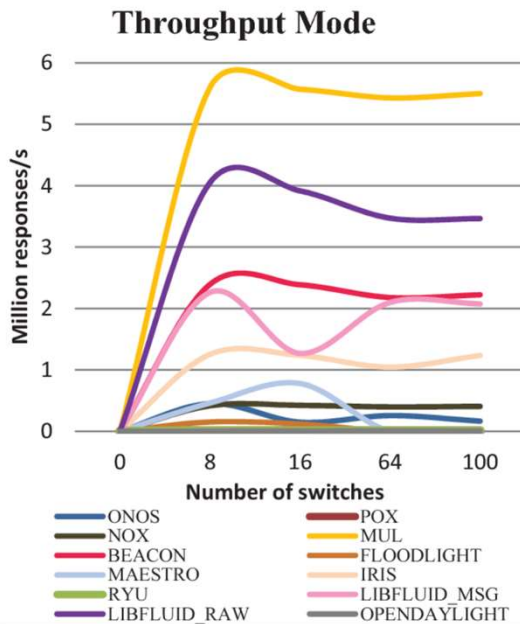


ACA ovs driver refactoring

Futurewei Cloud Lab

August 2021

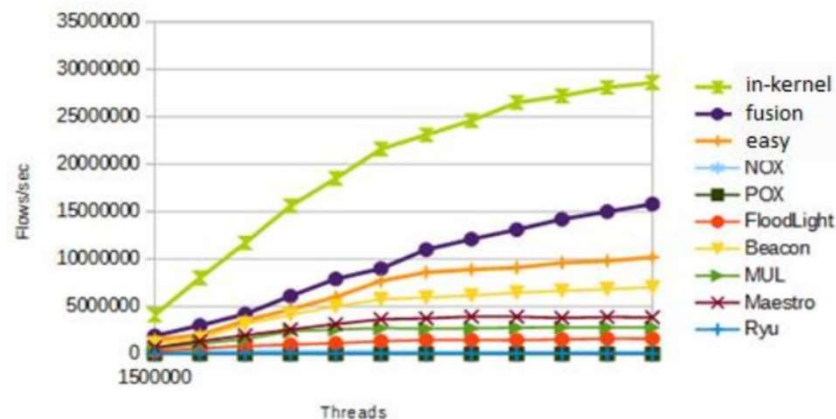
OpenFlow controller performance evaluation



- From multiple horizontal comparison reports/papers of popular OpenFlow controllers like left and more, we can summarize that the names with relatively good performance are:
 - MUL (C)
 - Libfluid_raw / Libfluid_msg (C++)
 - Beacon (java)
 - Maestro (java)
 - NOX_MultiThreading (C++)
- From past performance test, python ovs agent learning switch throughput is $\sim 1/70$ of libfluid, meaning
 - 0.03 million fps
 - Similar to RYU in left charts

OpenFlow controller performance evaluation

I/O throughput (cbench + l2learning), fps



- Performance is **10M fps** (based on libfluid).
- Latency is **55us**.

- From this report of an easy wrapper of Libfluid-msg (“easy” in left diagram), it beats Beacon, Maestro and others in
 - packet-in -> local L2 discovery lookup -> add-flow of neighbor
 - “Threads” axis is from 1 – 12
 - could reach 10M fps with 12 threads

OpenFlow controller performance evaluation

Switches: 16

Controllers: Beacon, NOX MT, Floodlight, raw, msg


Threads: 8

Application: learning switch

Command: `cbench -c localhost -p 6653 -m 10000 -l 16 -s 16 -M 1000000 -t`

Results (higher is better):

| Controller | flows/ms |
|------------|----------|
| raw | 7929.39 |
| Beacon | 7682.54 |
| msg | 5013.35 |
| Floodlight | 2490.61 |



OpenFlow controller performance evaluation

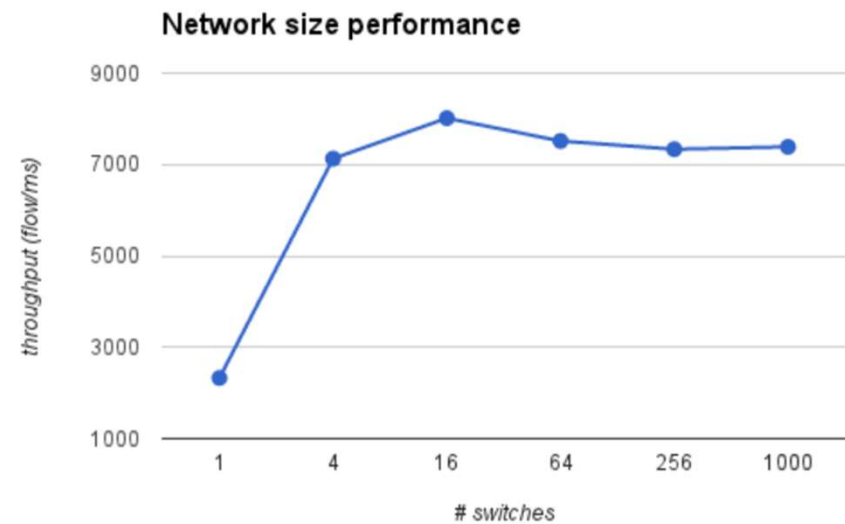
Threads: 8

Application: learning switch

Command: `cbench -c localhost -p 6653 -m 10000 -l 16 -s 16 -M 100000 -t`

Results (higher is better):

| Switches | flows/ms |
|----------|----------|
| 1 | 2327.88 |
| 4 | 7128.02 |
| 16 | 8013.98 |
| 64 | 7513.87 |
| 256 | 7334.20 |
| 1000 | 7386.28 |



OpenFlow controller design evaluation

- From these OpenFlow controller projects, we can learn from them in order to achieve higher performance
 - Efficient multi-threading non-blocking I/O model
 - Lower-level programming language
- Unlike them
 - We only focus on one type of OpenFlow switch – ovs currently, instead of raw OF protocol we need to understand and tune ovs OF behavior
 - We don't care much about number of switches (since we are not building a central controller), the maximum for us will be 4 (br-tun + br-int + ovsdb + aux connection for further enhancement later)
 - Thus, we value making the best use of single ovs connection the most

ACA V2 design

(including ovs-driver V2 refactoring)

- Layers
 - Upper
 - User update notification
 - Depending on upstream stack, can be grpc c/s or message queue library
 - Mid
 - Abstract goal states -> specific forwarding plane programming
 - Where ACA orchestrates and maintains states
 - Lower
 - Communicate with forwarding module, in our case is ovs (including ovs-db)
 - Where ovs-driver V2 refactoring takes place

ACA

Notification

grpc
client/server

MQ client

Goal-state

Goal-state
dispatcher

State programming orchestration

Data-plane programming

routes

acl

neighbors

...

dhcp

RA

arp

Configuration programming

port
binding

tunnel

Ovs driver

flows

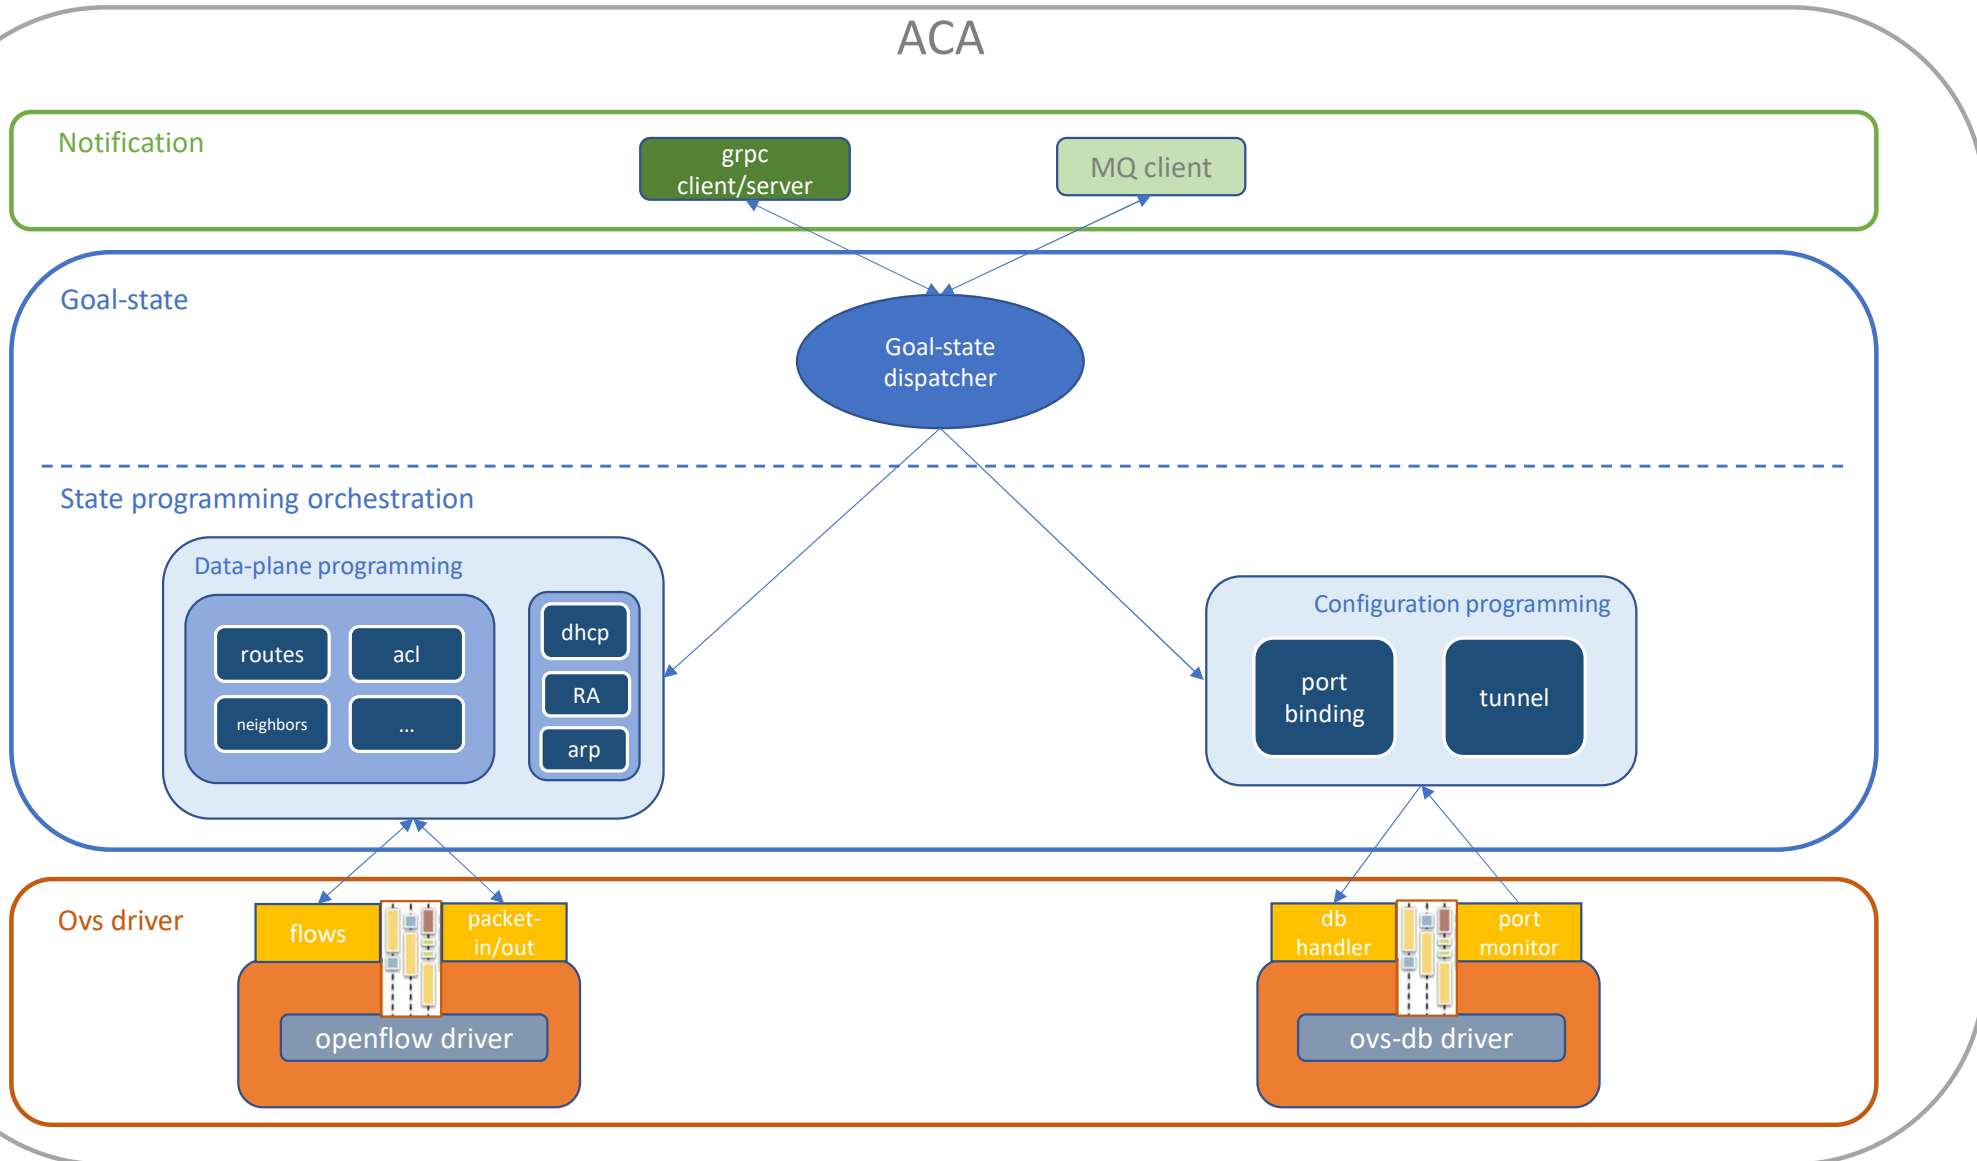
packet-
in/out

openflow driver


db
handler

port
monitor

ovs-db driver



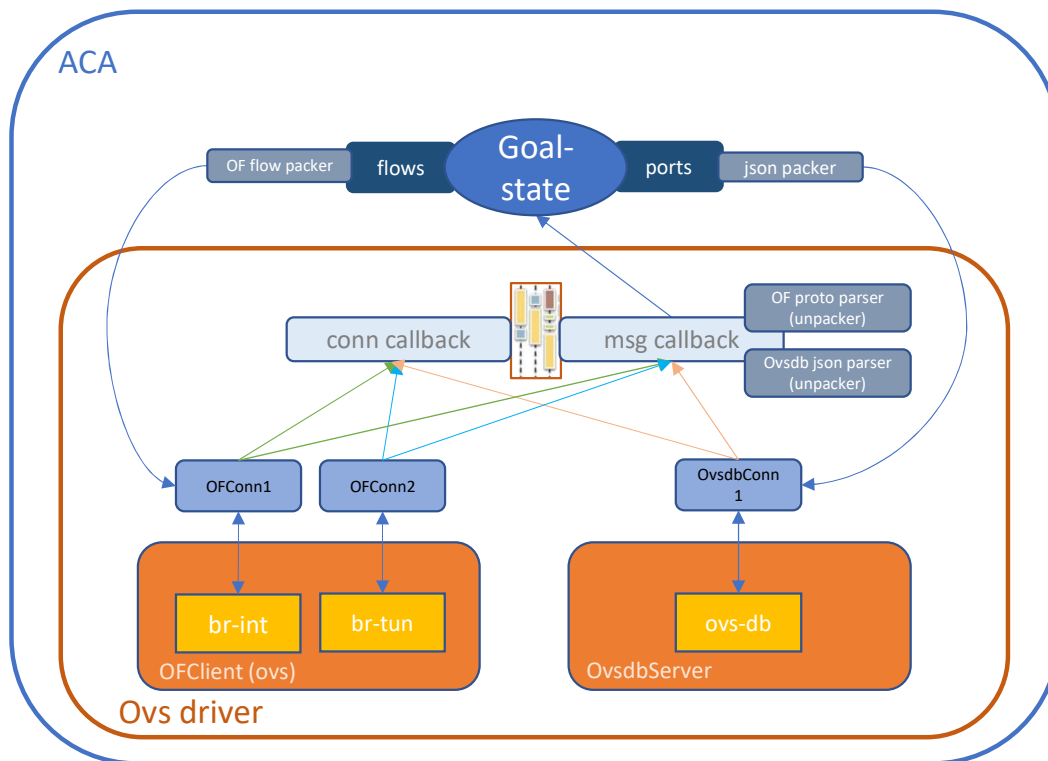
Libfluid fundamental



- Functions
 - Connection management
 - Event driven callbacks (non-blocking pipeline)
- Support numbers of connections (switches)
- High-performance server level event handling
- Based on raw OpenFlow protocol, needs some customization to support Ovs
- Does not include Ovsdb protocol/handler

ACA ovs-driver V2 design

break-down



- **Connections**

- br-int
- br-tun
- ovs-db

- **Connection callbacks**

- connection up
- connection down

- **Message callbacks (receiver/unpacker)**

- **openflow parser (from OFConn)**
 - Ovs reply (dump-flows, call reply with xid, bundle transaction reply etc.)
 - Packet-in
- **json parser (from OvsdbConn)**
 - db record change notifications (for example ports)
 - db query
 - transaction reply

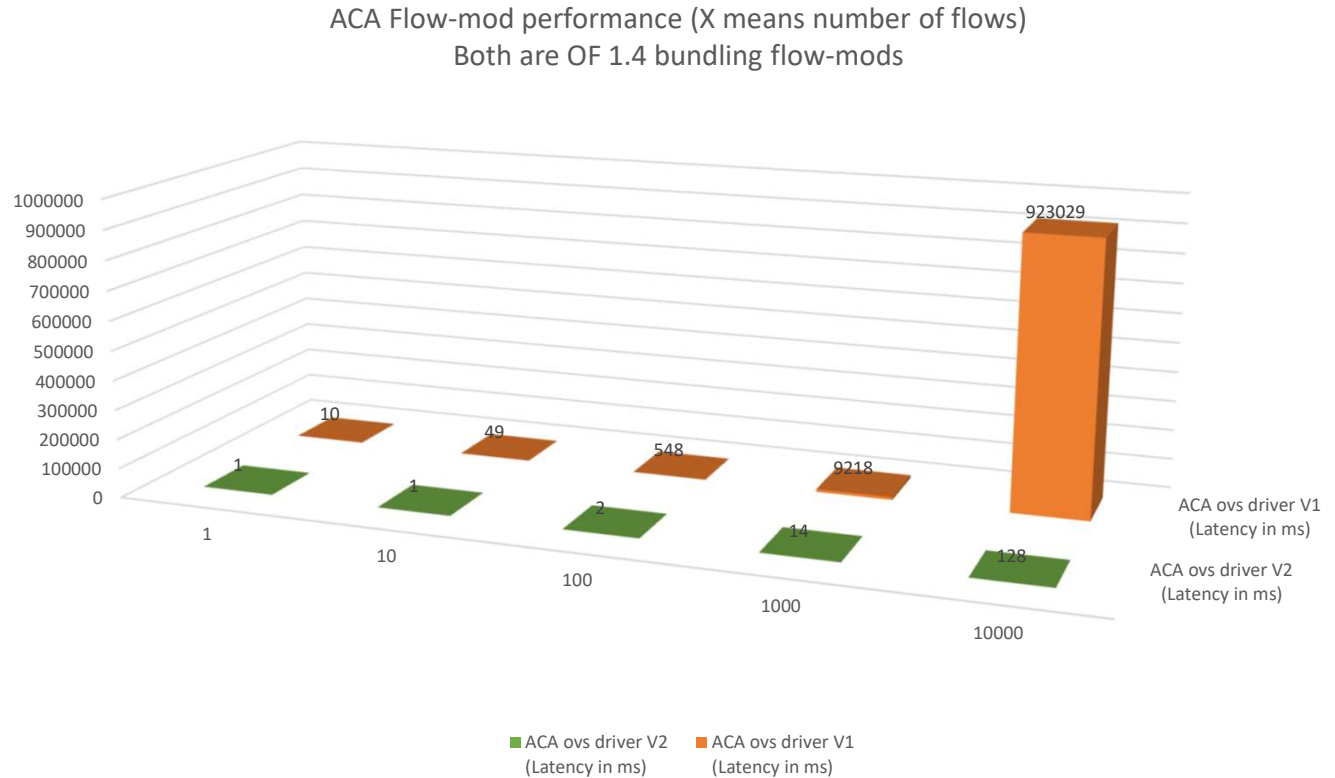
- **Programmer (sender/packer)**

- **OFConnection**
 - Construct flows -> add/mod/del flows (OF 1.3 and below)
 - Bundling flow-mods (OF 1.4 and above)
 - Packet out
- **OvsdbConnection**
 - Construct json rpc params -> ovsdb transactions

ACA ovs-driver V2

Performance improvement


| Flows | ACA ovs driver V2 (Latency in ms) | ACA ovs driver V1 (Latency in ms) |
|-------|--------------------------------------|--------------------------------------|
| 1 | 1 | 10 |
| 10 | 1 | 49 |
| 100 | 2 | 548 |
| 1000 | 14 | 9,218 |
| 10000 | 128 | 923,029 |



ACA ovs-driver V2

Performance expectation

| Flows | Bundling (OF1.4) Latency in ms | Non-bundling (OF1.3) Latency in ms |
|-------|-----------------------------------|--|
| 10 | 1 | 1 |
| 20 | 1 | 1 |
| 50 | 2 | 1 |
| 100 | 2 | 2 |
| 200 | 3 | 4 |
| 500 | 8 | 48 |
| 1k | 14 | 15 |
| 2k | 26 | 28 |
| 5k | 63 | 68 |
| 10k | 128 | 132 |
| 20k | 240 | 254 |
| 30k | 365 | 374 |
| 40k | 486 | 496 |
| 50k | 612 | 626 |



ACA ovs-driver V2

Conclusions of performance improvements

| Methods | Effect (measured by percentage of latency it saved) | Comments |
|---|---|--|
| Multiplexing non-blocking I/O model between ACA and ovs | 99 – 99.8% | <p>Instead of using V1 ofputil vconn which opens and closes each time and update flows sequentially, leverage non-blocking I/O model.</p> <p>Reduced add flow latency</p> <ul style="list-style-type: none">• For 100 rules from 548ms to 2ms which is 99.7%• For 1k rules from 9218ms to 14ms which is 99.8% |
| Use OF 1.4 bundling flow-mod transaction | <p>14 - 42% (if upgrade from OF 1.0)</p> <p>5 – 8% (if upgrade from OF 1.3)</p> | <p>From OF 1.0, reduced add flow latency</p> <ul style="list-style-type: none">• For 5k rules, from 110ms to 63ms which is 42%• For 15k rules, from 227ms to 195ms which is 14% <p>From OF 1.3, reduced add flow latency</p> <ul style="list-style-type: none">• For 5k rules, from 68ms to 63ms which is 8%• For 20k rules, from 254ms to 240ms which is 5.8% |